AI, Inequality, and Us

The below is entirely written by me, Ilan Rosenbaum, with no editing, advice, or writing done by any LLM or algorithmic help beyond spellcheck & the algorithms involved in the search tools I use for the research that went into this article.

AI, specifically Large Language Models (LLMs) are easily the most talked about anything in the country right now. In a country where the top 1% own 41.2% [1] of all the businesses. Those same businesses are giving us a worse deal than any time in recent history; the average markup on products has skyrocketed from 20% in the 1980s to 55% by 2020 [2]. It is with this backdrop that I'd like to approach the discussion of what AI is, what it can do, and its positioning in our cultural zeitgeist.

LLMs in white collar work is today's Chinese manufacturing boom, but with adoption moving at a substantially higher clip. Remember, accepting Chinese goods lead to the loss of domestic manufacturing and the expertise that made it possible. Replacing white-collar jobs and forcing reliance on AI to make up for lost productivity will lead to a similar loss of expertise and ability. This is outsourcing, it has the same benefits and downsides as it always does. This willing transfer of power is a classic Machiavellian misstep. ***If you need them, but they do not need you, or at their mercy be you the U.S. military, a tech giant, or a tiny startup.***

In order to make something with AI, you have to be able to describe it. As anyone who works in software knows, the customer can not accurately describe the full extent of what they need. If your company builds software your customers can easily describe, you're hosed. If what you make is truly complex, difficult to describe in its totality, and requires many decisions that have no objectively correct easily defined answer, you're going to be ok. Certain tasks, such as large migrations between products, languages, and frameworks have gotten substantially easier. But deciding on data structures, building roadmaps, and deciding on new features based all take just as long as it used to. In a world where switching is easier & new feature buildout is faster, but product & architecture decisions are as difficult as ever, the highest quality products will win more converts. For any good software team making complex and useful products, this is a growth opportunity, not an existential threat.

You need to be as good or better than the model is at whatever you are asking to get the best out of it. Regardless of the task — math, science, or software engineering — if you are not as skilled as the model, what it produces will be of unknown value to you. Why is that? Because model outputs are jagged. They may have extensive knowledge in a specific aspect of a subject but almost none in a seemingly similar, neighboring idea, or the same exact idea they knew a ton about prompted a second time. See [3] for more details on the nature of model jaggedness. This is one cause of something all regular LLM users have seen hundreds of times: blatant and confident lies. This fact rules out usage of AI without oversight in any realm where one wrong answer is unacceptable. Or at least it should. Important examples of this would be, but are not limited to, healthcare, war, security, and law. These are realms AI is already penetrating. Jaggedness should be unacceptable in the drawing up kill lists. It leads to citing cases that do not exist in legal motions. This is not just a political issue, if you use AI to do something you personally don't understand and it spits out plausible looking garbage, then you present it as your own work, you will look like a fool. Worse than looking like a fool, it will show the people you work with you are the person AI should replace. Would you hire the lawyer that submitted AI generated motions on your behalf? **If you are no better than a ChatGPT wrapper, what value do you provide?**

Today's LLMs are many things, but to me, first and foremost, they are incredibly powerful search engines. They have the ability to synthesize and repeat novel combinations of more or less everything humans have written and maintained. This was made possible by AI companies taking the sum of human knowledge as it existed on the internet — almost all of which they did not have the author(s) permission to take [16] [17] [18] [19] — and running it through algorithms that then store a complex and impossible to fully decipher version of that information within the confines of massive matrices. I am not an expert on the 'how'. If you want a better understanding I'd recommend 3 Brown 1 Blue's series on Deep Learning. What's important to understand is these models are not made: nobody decides exactly how they choose to spit out what they spit out, instead we grow them by punishing wrong answers and rewarding correct ones. When trying to synthesize general intelligence, some incredibly important questions come up. What is a right answer? For math, coding, and some of the sciences this is simple: the right answer is the one you can prove. But what about the arts, or any question without a definitive, provable answer? Well I guess the right answer is some average of all the text humans have ever written, guided of course in the direction corporations find most palatable. Does that mean

directly spitting out the words of famous authors? Repeating the opinions of Redditors, LinkedIn warriors, 4Chan users?? What about the increasing amount of the internet that is itself AI generated. That is an open question. Whatever they decide is the 'correct' answer that these models will be emulating is going to shape writing, culture, and human thought in profound and unknowable ways.

Due to the fact AI is trained on human data, it is flawed in the same ways as the data and can systematize bias against minorities. Sorry, not it can, it has, "researchers found that race, gender, income, and housing status influenced evaluation and treatment recommendations — even when patients had the same health conditions" [7]. Again not only do these findings have horrific societal implications, but it also means that leaning too heavily on this will lead to worse care overall. If we harken back to the example of Chinese manufacturing, we may have more goods now but lots of it is way worse. This is ok in some industries: quick, almost free, and mediocre is ok for analyzing data small business owners didn't used to get any insights from. It's great for throwing together a prototype for an idea you have to see if it's got legs. But as a mass surveillance apparatus? Being ok with government surveillance is a stretch for many of us as is, but what if you know, for certain, that the surveillance was both biased and confidently wrong? Misuse won't just be cruel and vindictive, it would be arbitrary and pointless.

Finally, a point that can not be ignored is that these systems rely on massive data centers. The centers are worse for air quality than natural gas power plants [20]. But unlike a power plant, they increase your power bill [21]. This hits close to home for me as QTS was threatening to build a power plant near where I live — a plant that would have consumed more than twice the power of Dane county, home of over half a million people. Because of how quickly they are rushing to build out capacity, and the amount of money pouring into the buildout effort, the cost effectiveness and energy efficiency of these data centers doesn't seem to matter. This isn't just conjecture, "The carbon intensity of electricity used by data centers [is] 48% higher than the US average." per James O'Donnell and Casey Crownhart at the MIT Technology Review [22]. Sadly that is not where it ends. These things are so damn loud "You feel it in your bones" [23]. Constant noise is not simply annoying, it has real health effects. There is decades of research showing noise pollution causes chronic heart issue and heart attacks along with links to schizophrenia, anxiety, and depression [24]. There is plenty more from land and water use, to the fact these things are ugly as hell. I'm certain there is at least one aspect of AI data centers that

would make you hesitant before inviting one into your back yard. ***If you don't want your power bill to go up, your air to be polluted, or to literally have a heart attack you are in favor of slowing down this buildout so we can regulate how these behemoths are built.***

I am not advocating for the halting of AI development, this is a breakthrough to be sure, but let's build these data centers responsibly with high standards for energy efficiency, limited noise levels, and in a way that doesn't pass on the economic and health costs to the already financially stretched and unhealthy middle American.

We have the power to shape the direction of this industry. An industry claiming to have the capacity to further centralize power into the hands of the few who already have so much. If we regulate the buildout of data centers, there will be better infrastructure for us all. If we use, and *push our employees to use* AI responsibly, we will become more productive without sacrificing our current and future abilities as people and as corporations. If we disallow uses of LLM technology without oversight in fields where there truly is no room for error, we save ourselves from the dystopia that is Shein healthcare and Temu legal representation. To think the effects of AI will be as overwhelmingly positive as the gilded dreams of those whose careers rely on selling it, would be, well, hallucinogenic.

Now what?

Find out if a data center is being built near you. If one is, work with your local community to fight against it or to fight to make sure it's built in a way that helps you more than it hurts. I promise there are already groups locally that are trying to do this. Join them.

Don't blindly trust your representatives. In my city it was found those on the board who accepted the deal initially were in the pocket of the developers. Send emails to your state and federal representatives — I know it sounds silly and seems like it'd be a lot of work but almost nobody does it. Your voice will not be heard if you do not speak up. Want specific templates for reaching out to your representatives? Future of Life has a wonderful page dedicated to this: https://futureoflife.org/take-action/contact-your-legislator/. Tell the people in power what you

care about.

Maybe you worry about water that you think the benefits of AI in healthcare outweigh the costs. That's ok. Maybe you care about an issue I did not bring up and think I'm fear mongering about half the stuff I wrote. That's ok too — hell I managed not to bring up deepfakes once and they're being used to propagandize the world in a way never seen before. This isn't about exclusively solving the issues I care about, let alone the few I wrote about here. We need to make sure this new, transformational technology is regulated so that it helps us. That means speaking up about what matters to you. If you have the means, donate to groups in support of well written legislation. Democracy only works if we participate. In this day in age that means money, but it also means time.

If you have the authority to represent an organization, please read this and sign on behalf of said organization if you agree with it: https://humanstatement.org/

Finally, and most importantly, talk to your friends and co-workers. See if they are worried about this, share your viewpoints. Even just one person. No set of isolated actions solve this. Only together can we solve issues this big.

---

1. https://www.nber.org/papers/w31823
2. https://www.econstor.eu/handle/10419/235280
3. https://www.pnas.org/doi/epdf/10.1073/pnas.2400489121
4. https://arxiv.org/abs/2601.07573
5. https://www.aljazeera.com/news/2024/4/4/ai-assisted-genocide-israel-reportedly-used-database-for-9000-kill-lists
6. https://mashable.com/article/samsung-chatgpt-leak-details
7. https://codex.ucsf.edu/news/editors-pick-study-finds-ai-medical-tools-show-bias-potential-misdiagnosis-and-patient-harm
8. https://www.cbsnews.com/news/ai-executive-dario-amodei-on-the-end-linux-anthropic-would-not-cross/
9. https://mason.com/mlaab/2024/02/16/2000-sanction-in-another-ai-hallucinated-citation-case/
10. https://scholarscompass.vcu.edu/cgi/viewcontent.cgi?article=1101&context=law_pubs/
11. https://www.eesi.org/articles/view/data-center-power-demands-are-contributing-to-higher-energy-bills
12. https://www.technologyreview.com/2025/05/20/1116547/ai-energy-usage-climate-footprint-big-tech
13. https://water.usgs.gov/vtlab/water-availability/03-vulnerability
14. https://www.nyr.org/2025/08/27/ai-x1-5461467/ai-is-driving-a-data-center-boom-in-rural-america-locals-are-divided-on-the-benefits
15. https://www.youtube.com/watch?v=mJTUOW_ia
16. https://people.cs.vt.edu/~clifver/assets/papers/memo12.pdf
17. https://en.wikipedia.org/wiki/Authors_Guild,_Inc._v._Google,_Inc.
18. https://antechnica.com/tech-policy/2025/02/meta-defends-its-vast-book-torrenting-were-just-a-leech-no-proof-of-seeding/
19. https://harmas.bearblog.dev/the-great-scrape/
20. https://abcchicago.com/post/what-is-data-center-expect-more-centers-illinois-like-cyngenus-aurora-4-amid-noise-electric-bill-complaints/28810141/
21. https://cseh.sulaxis.edu/research/how-noise-pollution-quietly-affects-your-health